

# Source\_Assignment\_Manola

July 9, 2023

```
[1]: print ("The steps that I will follow will be the following:")
      print ("1. Introduction")
      print ("2. Exploration of Indoors temperature dataset")
      print ("3. Stationarity Tests")
      print ("4. Univariate Temperature forecasting using smoothing")
      print ("5. Univariate Temperature forecasting using SARIMA Model")
      print ("6. Multivariate Temperature forecasting using SARIMAX Model")
      print ("7. Final Comparison of Models Using MSE")
```

The steps that I will follow will be the following:

1. Introduction
2. Exploration of Indoors temperature dataset
3. Stationarity Tests
4. Univariate Temperature forecasting using smoothing
5. Univariate Temperature forecasting using SARIMA Model
6. Multivariate Temperature forecasting using SARIMAX Model
7. Final Comparison of Models Using MSE

```
[2]: from IPython.display import Markdown

      comment = "**1. Introduction**"
      font_size = "20px"

      Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</
      ↪span>")
```

[2]: 1. Introduction

```
[3]: print("Our goal for this assignment is to predict the indoors temperature of a
      ↪greenhouse for every hour for the \
      next 24 hours. To start with this, let's have a look at our available data in
      ↪order to draw a good method for our \
      forecasting methods. In the ReadMe.pdf several tables are given, but not all of
      ↪them are relevant for the indoors \
      greenhouse temperature. For example the 'Production' and 'Crop parameters' are
      ↪relevant for the growth and yield \
      of the plants, but not for the temperature of the greenhouse. We won't use
      ↪these tables. We are going to load the \
```

```
two relevant tables that are the 'Weather data', that gives information of the
↳weather conditions outside of the \
greenhouse, and the 'Greenhouse climate' that gives information about the
↳indoor climate, the status of actuators \
and irrigation.")
```

Our goal for this assignment is to predict the indoors temperature of a greenhouse for every hour for the next 24 hours. To start with this, let's have a look at our available data in order to draw a good method for our forecasting methods. In the ReadMe.pdf several tables are given, but not all of them are relevant for the indoors greenhouse temperature. For example the 'Production' and 'Crop parameters' are relevant for the growth and yield of the plants, but not for the temperature of the greenhouse. We won't use these tables. We are going to load the two relevant tables that are the 'Weather data', that gives information of the weather conditions outside of the greenhouse, and the 'Greenhouse climate' that gives information about the indoor climate, the status of actuators and irrigation.

```
[4]: print("let's import some useful libraries first")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from datetime import datetime
from datetime import timedelta
from dateutil.relativedelta import relativedelta
from IPython.display import display
import os
warnings.filterwarnings('ignore')
pd.options.display.float_format = '{:,.5f}'.format
%matplotlib inline
plotsize = (13, 5)
```

let's import some useful libraries first

```
[5]: print("Let's have a first look at the table of the outdoors weather conditions.
↳")
df_outd = pd.read_csv("/Users/irismanola/Documents/ML/Source_Assignment/Weather.
↳csv")
df_outd.columns
df_outd.head()
```

Let's have a first look at the table of the outdoors weather conditions.

```
[5]:
```

	time	AbsHumOut	Iglob	PARout	Pyrgeo	RadSum	Rain	\
0	43,815.00000	6.22095	0.00000	0.00000	-72.00000	215.00000	0.00000	
1	43,815.00347	6.22095	0.00000	0.00000	-73.00000	0.00000	0.00000	
2	43,815.00694	6.20556	0.00000	0.00000	-76.00000	0.00000	0.00000	
3	43,815.01042	6.19017	0.00000	0.00000	-77.00000	0.00000	0.00000	
4	43,815.01389	6.16262	0.00000	0.00000	-75.00000	0.00000	0.00000	

	Rhout	Tout	Winddir	Windsp
0	80.60000	6.90000	32.00000	4.70000
1	80.60000	6.90000	32.00000	4.70000
2	80.40000	6.90000	32.00000	4.70000
3	80.20000	6.90000	32.00000	4.70000
4	80.90000	6.70000	32.00000	4.70000

```
[6]: print("And now let's have a look at the indoors conditions.")

df_ind = pd.read_csv("/Users/irismanola/Documents/ML/Source_Assignment/
↳GreenhouseClimate.csv")
df_ind.columns
df_ind.head()
```

And now let's have a look at the indoors conditions.

```
[6]:
```

	time	AssimLight	BlackScr	CO2air	Cum_irr	EC_drain_PC	\
0	43,815.00000	0.00000	35.00000	472.00000	0.35480	3.98000	
1	43,815.00347	0.00000	85.00000	501.00000	0.53220	3.99000	
2	43,815.00694	0.00000	96.00000	489.00000	0.53220	3.99000	
3	43,815.01042	0.00000	96.00000	497.00000	0.53220	4.00000	
4	43,815.01389	0.00000	96.00000	477.00000	0.53220	4.00000	

	EnScr	HumDef	PipeGrow	PipeLow	...	t_rail_min_sp	t_rail_min_vip	\
0	100.00000	7.17000	0.00000	45.20000	...	NaN	0.00000	
1	100.00000	6.94000	0.00000	43.60000	...	NaN	0.00000	
2	100.00000	7.24000	0.00000	42.30000	...	NaN	0.00000	
3	100.00000	6.74000	0.00000	41.20000	...	NaN	0.00000	
4	100.00000	6.71000	0.00000	41.20000	...	NaN	0.00000	

	t_vent_sp	t_ventlee_vip	t_ventwind_vip	water_sup	\
0	NaN	25.00000	26.00000	4.00000	
1	NaN	25.00000	26.00000	6.00000	
2	NaN	25.00000	26.00000	6.00000	
3	NaN	25.00000	26.00000	6.00000	
4	NaN	25.00000	26.00000	6.00000	

	water_sup_intervals_sp_min	water_sup_intervals_vip_min	window_pos_lee_sp	\
0	NaN	1,000.00000	NaN	
1	NaN	1,000.00000	NaN	

2	NaN	1,000.00000	NaN
3	NaN	1,000.00000	NaN
4	NaN	1,000.00000	NaN

window_pos_lee_vip	
0	1.20000
1	1.20000
2	1.20000
3	1.20000
4	1.20000

[5 rows x 50 columns]

```
[7]: print("the columns are too many, so we do the following to see them all:")
pd.set_option('display.max_columns', None)
df_ind.head()
```

the columns are too many, so we do the following to see them all:

```
[7]:
```

	time	AssimLight	BlackScr	CO2air	Cum_irr	EC_drain_PC	\
0	43,815.00000	0.00000	35.00000	472.00000	0.35480	3.98000	
1	43,815.00347	0.00000	85.00000	501.00000	0.53220	3.99000	
2	43,815.00694	0.00000	96.00000	489.00000	0.53220	3.99000	
3	43,815.01042	0.00000	96.00000	497.00000	0.53220	4.00000	
4	43,815.01389	0.00000	96.00000	477.00000	0.53220	4.00000	

	EnScr	HumDef	PipeGrow	PipeLow	Rhair	Tair	Tot_PAR	\
0	100.00000	7.17000	0.00000	45.20000	60.60000	20.90000	0.00000	
1	100.00000	6.94000	0.00000	43.60000	61.40000	20.70000	0.00000	
2	100.00000	7.24000	0.00000	42.30000	60.90000	21.20000	0.00000	
3	100.00000	6.74000	0.00000	41.20000	62.50000	20.70000	0.00000	
4	100.00000	6.71000	0.00000	41.20000	62.70000	20.70000	0.00000	

	Tot_PAR_Lamps	VentLee	Ventwind	assim_sp	assim_vip	co2_dos	co2_sp	\
0	0.00000	0.80000	0.00000	NaN	0.00000	NaN	NaN	
1	0.00000	0.80000	0.00000	NaN	0.00000	0.00034	NaN	
2	0.00000	0.80000	0.00000	NaN	0.00000	0.00036	NaN	
3	0.00000	0.80000	0.00000	NaN	0.00000	0.00201	NaN	
4	0.00000	0.80000	0.00000	NaN	0.00000	0.00155	NaN	

	co2_vip	dx_sp	dx_vip	int_blue_sp	int_blue_vip	int_farred_sp	\
0	400.00000	NaN	2.20000	NaN	NaN	NaN	
1	400.00000	NaN	2.20000	0.00000	NaN	0.00000	
2	400.00000	NaN	2.20000	0.00000	NaN	0.00000	
3	400.00000	NaN	2.20000	0.00000	NaN	0.00000	
4	400.00000	NaN	2.20000	0.00000	NaN	0.00000	

	int_farred_vip	int_red_sp	int_red_vip	int_white_sp	int_white_vip	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	0.00000	NaN	0.00000	NaN	
2	NaN	0.00000	NaN	0.00000	NaN	
3	NaN	0.00000	NaN	0.00000	NaN	
4	NaN	0.00000	NaN	0.00000	NaN	

	pH_drain_PC	scr_blk_sp	scr_blk_vip	scr_enrg_sp	scr_enrg_vip	\
0	6.27000	NaN	96.00000	NaN	100.00000	
1	6.28000	NaN	96.00000	NaN	100.00000	
2	6.28000	NaN	96.00000	NaN	100.00000	
3	6.27000	NaN	96.00000	NaN	100.00000	
4	6.28000	NaN	96.00000	NaN	100.00000	

	t_grow_min_sp	t_grow_min_vip	t_heat_sp	t_heat_vip	t_rail_min_sp	\
0	NaN	10.00000	NaN	21.00000	NaN	
1	NaN	10.00000	NaN	21.00000	NaN	
2	NaN	10.00000	NaN	21.00000	NaN	
3	NaN	10.00000	NaN	21.00000	NaN	
4	NaN	10.00000	NaN	21.00000	NaN	

	t_rail_min_vip	t_vent_sp	t_ventlee_vip	t_ventwind_vip	water_sup	\
0	0.00000	NaN	25.00000	26.00000	4.00000	
1	0.00000	NaN	25.00000	26.00000	6.00000	
2	0.00000	NaN	25.00000	26.00000	6.00000	
3	0.00000	NaN	25.00000	26.00000	6.00000	
4	0.00000	NaN	25.00000	26.00000	6.00000	

	water_sup_intervals_sp_min	water_sup_intervals_vip_min	window_pos_lee_sp	\
0	NaN	1,000.00000	NaN	
1	NaN	1,000.00000	NaN	
2	NaN	1,000.00000	NaN	
3	NaN	1,000.00000	NaN	
4	NaN	1,000.00000	NaN	

	window_pos_lee_vip
0	1.20000
1	1.20000
2	1.20000
3	1.20000
4	1.20000

```
[8]: print(df_outd.shape)
```

```
(47809, 11)
```

```
[9]: print(df_ind.shape)
```

(47809, 50)

```
[10]: print(f"Both tables have the same lengths. Every row is a measurement every 5_
      ↪minutes. Then have have data for \
      {47809*5/60:.1f} hours, which is {47809 * 5/60/24:.1f} days of data, or_
      ↪approximately {47809 * 5/60/24/30:.1f} months.")
```

Both tables have the same lengths. Every row is a measurement every 5 minutes. Then have have data for 3984.1 hours, which is 166.0 days of data, or approximately 5.5 months.

```
[11]: print ("In the ReadMe.pdf all the variable names are explained. The Tair_
      ↪(Greenhouse Air temperature) is the \
      variable that we have to forecast. To do so we can use either univariate_
      ↪(involving only the Tair timeseries) \
      or multivariate forecasting methods (involving other given variables that_
      ↪influence the Tair). \
      We will start with simple univariate methods and will move towards more complex_
      ↪multivariate methods. For now let's \
      focus only on the Tair and explore it.")
```

In the ReadMe.pdf all the variable names are explained. The Tair (Greenhouse Air temperature) is the variable that we have to forecast. To do so we can use either univariate (involving only the Tair timeseries) or multivariate forecasting methods (involving other given variables that influence the Tair). We will start with simple univariate methods and will move towards more complex multivariate methods. For now let's focus only on the Tair and explore it.

```
[12]: from IPython.display import Markdown

      comment = "**2. Exploration of Indoors temperature dataset**"
      font_size = "20px"

      Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</_
      ↪span>")
```


[12]: 2. Exploration of Indoors temperature dataset

```
[13]: print ("Let's see if the time is well set every 5 minutes and without missing_
      ↪variables.")
      df_ind[['time']].plot(figsize=plotsize)

      plt.title('Time')
```

Let's see if the time is well set every 5 minutes and without missing variables.

```
[13]: Text(0.5, 1.0, 'Time')
```



```
[14]: print ("are there any missing values in the time ?")
time=df_ind[['time']]
time.isnull().sum()
```

are there any missing values in the time ?

```
[14]: time    0
      dtype: int64
```

```
[15]: print ("the time variable looks nice and linear as it should and there are no
      ↪missing variables.")
```

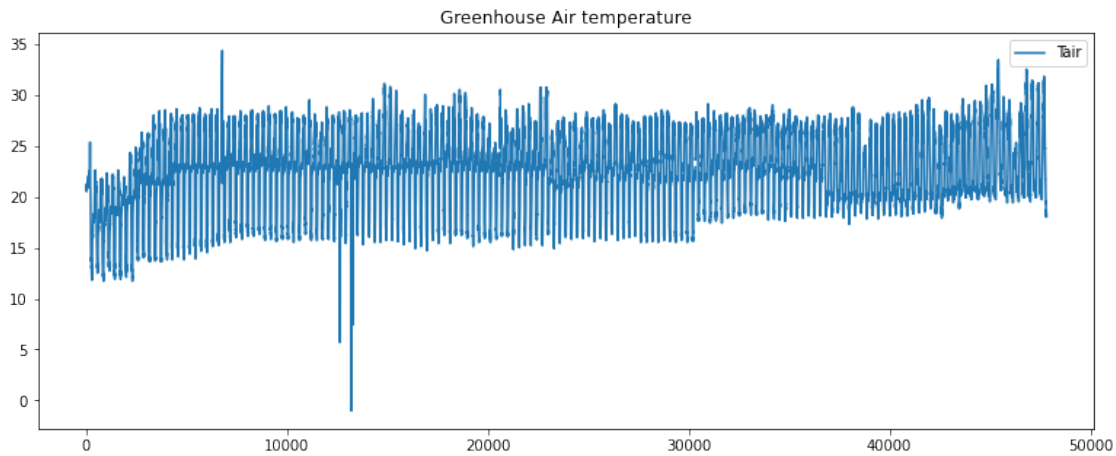
the time variable looks nice and linear as it should and there are no missing variables.

```
[16]: Tair = df_ind[['Tair']]
      Tair.head(10)
```

```
[16]:      Tair
0  20.90000
1  20.70000
2  21.20000
3  20.70000
4  20.70000
5  20.60000
6  20.70000
7  20.60000
8  20.70000
9  21.00000
```

```
[17]: Tair.plot(figsize=plotsize)
      plt.title('Greenhouse Air temperature')
```

```
[17]: Text(0.5, 1.0, 'Greenhouse Air temperature')
```



```
[18]: statistics = Tair.describe()
print(statistics)
```

	Tair
count	47,738.00000
mean	22.71351
std	3.92040
min	-1.00000
25%	20.00000
50%	22.80000
75%	26.40000
max	34.32632

```
[19]: print ("There are no crazy outliers to indicate a large mistake in the
    ↳ measurements that need to be removed. \
There is a warm outlier at the beginning of the measurments and two cold ones. \
    ↳ The cold ones might be cold \
nights were the window of the greenhouse were left (forgotten?) open. \
There is no heretoscedasticity in the data and no obvious strong \
    ↳ autocorrelation. We can see some trend at the beginning \
of the data, probably indicating a seasonal change, like transition from spring \
    ↳ to summer. There is also a more \
abrupt temperature drop after the second half of the data, probably another \
    ↳ change in season from summer to autumn. \
We also see a large periodicity that looks like an intense daily cicle. Let's \
    ↳ zoom in to see if this is indeed \
a daily cycle.")
print (" ")
```



```
print ("Further below we will see why some of my preliminary assumptions are_\n\n→not correct!")
```

There are no crazy outliers to indicate a large mistake in the measurements that need to be removed. There is a warm outlier at the beginning of the measurements and two cold ones. The cold ones might be cold nights where the window of the greenhouse were left (forgotten?) open. There is no heteroscedasticity in the data and no obvious strong autocorrelation. We can see some trend at the beginning of the data, probably indicating a seasonal change, like transition from spring to summer. There is also a more abrupt temperature drop after the second half of the data, probably another change in season from summer to autumn. We also see a large periodicity that looks like an intense daily cycle. Let's zoom in to see if this is indeed a daily cycle.


Further below we will see why some of my preliminary assumptions are not correct!

```
[20]: start_index = 10000
      end_index = 15000

      Tair_slice = Tair[start_index:end_index]

      Tair_slice.plot(figsize=plotsize)
      plt.title('Greenhouse Air temperature, zoomed in')
```

```
[20]: Text(0.5, 1.0, 'Greenhouse Air temperature, zoomed in')
```



```
[21]: print ("Yes, this looks indeed like a daily cycle (there are 288 - 5 minutes in_\n\n→a day). The temperature change \n\nis not Gaussian though. We see an evening temperature drop, \n\n")
```

```
then in the morning a sharp temperature increase, a plateau and a second_
↳temperature increase. This needs some \
thinking. Let's come back to this later again.")
```

Yes, this looks indeed like a daily cycle (there are 288 - 5 minutes in a day). The temperature change is not Gaussian though. We see an evening temperature drop, then in the morning a sharp temperature increase, a plateau and a second temperature increase. This needs some thinking. Let's come back to this later again.

```
[22]: Tair.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47809 entries, 0 to 47808
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Tair    47738 non-null    float64
dtypes: float64(1)
memory usage: 373.6 KB
```

```
[23]: Tair.isnull().sum()
```

```
[23]: Tair    71
      dtype: int64
```

```
[24]: print("We see that there are 71 missing values among the 47.809 numerical_
      ↳entries.")
      print ("Let's fill the missing values. They are not many and they are only_
      ↳every 5 minutes, so the method we \
      chose won't make a big difference. Let's do a simple forward fill.")
```

We see that there are 71 missing values among the 47.809 numerical entries. Let's fill the missing values. They are not many and they are only every 5 minutes, so the method we chose won't make a big difference. Let's do a simple forward fill.

```
[25]: Tair = Tair.fillna(method='ffill')
      Tair.isnull().sum()
```

```
[25]: Tair    0
      dtype: int64
```

```
[26]: print ("cool, there are no missing values any more.")
      print ("Since we are interested in hourly observations let's smoothen the data_
      ↳from 5-mins to hourly.")
```

```

# Because the length of Tair (47809) is not divided exactly with 12 I make the
↳ number of groups and a new length to fit them
num_groups = len(Tair) // 12
new_length = num_groups * 12

# Reshape the DataFrame with 12 rows
Tair_reshape = Tair.values[:new_length].reshape(num_groups, 12)

# Calculate the mean along the rows
Tair_mean = np.mean(Tair_reshape, axis=1)

# Create a new DataFrame with the averaged values
Tair_hr = pd.DataFrame(Tair_mean, columns=['Value'])


Tair_hr.plot(figsize=plotsize)
plt.title('Greenhouse Air temperature, hourly ')

```

cool, there are no missing values any more.

Since we are interested in hourly observations let's smoothen the data from 5-mins to hourly.

```
[26]: Text(0.5, 1.0, 'Greenhouse Air temperature, hourly ')
```



```
[27]: Tair_hr.head()
```

```

[27]:      Value
0  20.77500
1  20.74167
2  20.92500
3  21.00833
4  21.11667

```

```
[28]: len(Tair_hr)
```

```
[28]: 3984
```

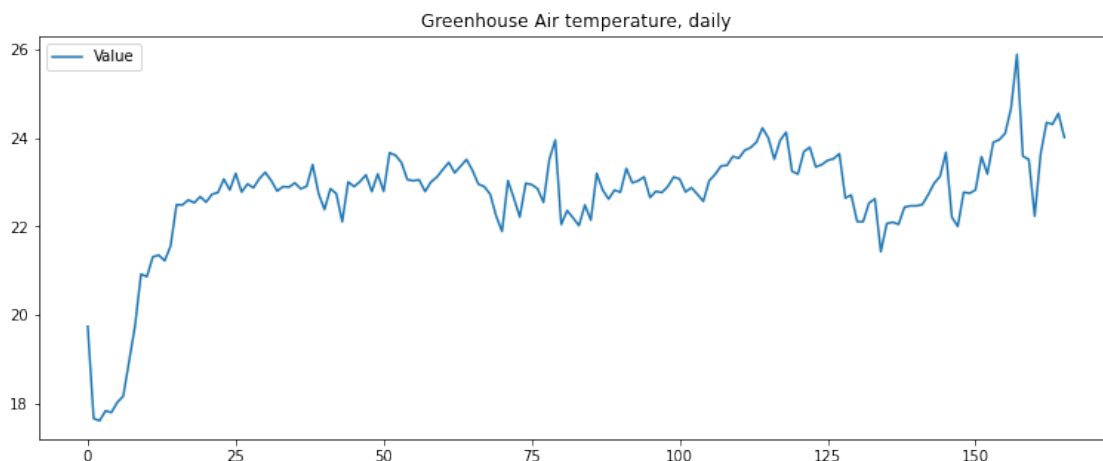
```
[29]: print ("The graph still looks quite noisy. Let's plot the daily values as well.
↪")

num_groups = len(Tair) // 288
new_length = num_groups * 288
Tair_reshape = Tair.values[:new_length].reshape(num_groups, 288)
Tair_mean = np.mean(Tair_reshape, axis=1)
Tair_day = pd.DataFrame(Tair_mean, columns=['Value'])

Tair_day.plot(figsize=plotsize)
plt.title('Greenhouse Air temperature, daily')
```

The graph still looks quite noisy. Let's plot the daily values as well.

```
[29]: Text(0.5, 1.0, 'Greenhouse Air temperature, daily')
```



```
[30]: print ("This is interesting! We now clearly see the trend at the beginning (due
↪to change of season most likely). \
We also see some autocorrelation that was not visible in the 5-mins and hourly
↪data because of the noise.")
```

This is interesting! We now clearly see the trend at the beginning (due to change of season most likely). We also see some autocorrelation that was not visible in the 5-mins and hourly data because of the noise.

```
[31]: print ("Out of curiosity, let's see how the temperature inside the greenhouse
↪relates to the observed temperature \
```

```
outside of it.")
```


Out of curiosity, let's see how the temperature inside the greenhouse relates to the observed temperature outside of it.

```
[32]: Tair_out = df_outd[['Tout']]

fig, ax = plt.subplots(figsize=plotsize)
ax.plot(Tair_out.index, Tair_out, color='red', label='Tair_out')
ax.plot(Tair.index, Tair, color='blue', label='Tair')

ax.set_xlabel('Time')
ax.set_ylabel('Temperature')
ax.legend()
ax.set_title('Indoors and outdoors 5-min temperature')

plt.show()
```



```
[33]: Tair_out = df_outd[['Tout']]
num_groups = len(Tair_out) // 288
new_length = num_groups * 288
Tair_reshape = Tair_out.values[:new_length].reshape(num_groups, 288)
Tair_mean = np.mean(Tair_reshape, axis=1)
Tairout_day = pd.DataFrame(Tair_mean, columns=['Value'])


fig, ax = plt.subplots(figsize=plotsize)
ax.plot(Tairout_day.index, Tairout_day, color='red', label='Tair_out')
ax.plot(Tair_day.index, Tair_day, color='blue', label='Tair')
```

```

ax.set_xlabel('Time')
ax.set_ylabel('Temperature')
ax.legend()
ax.set_title('Indoors and outdoors daily temperature')

plt.show()

```



```

[34]: print ("We can see that the outdoors temperature timeseries is non-stationary,
        ↳as there is heteroscedasticity and \
        autocorrelation and seasonality.")
print ("")
print ("By comparing the in- with the out-door temperature we can see how more
        ↳smooth the indoor temperature is. \
        We see that the average difference in T in and out is more than 10 degrees. We
        ↳can also see that my initial assumption \
        that there is a transition from spring to summer to autumn was not correct! We
        ↳rather see a transition from a \
        (cold) spring to more summer-like temperatures (assuming this is data from the
        ↳Netherlands of course!). My assumption \
        now is that the increase that we see in the T in at the beginning of the
        ↳measurments and the abrupt drop after the \
        second half is that this is intentionally done to help the plant growth.")

```

We can see that the outdoors temperature timeseries is non-stationary, as there is heteroscedasticity and autocorrelation and seasonality.

By comparing the in- with the out-door temperature we can see how more smooth the indoor temperature is. We see that the average difference in T in and out is more than 10 degrees. We can also see that my initial assumption that there is a transition from spring to summer to autumn was not correct! We rather see a

transition from a (cold) spring to more summer-like temperatures (assuming this is data from the Netherlands of course!). My assumption now is that the increase that we see in the T in at the beginning of the measurements and the abrupt drop after the second half is that this is intentionally done to help the plant growth.

```
[35]: comment = "***3. Stationarity Tests**"
font_size = "20px"

Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</span>")
```

[35]: 3. Stationarity Tests

```
[36]: print ("It is important to determine whether the Tair is stationary. This will \
        ↳determine which methods we will \
        use to forecast it, and will tell whether we have to transform the data to get \
        ↳a cleaner signal (in non-stationary).")
print ("In order for a timeseries data to be stationary, the data must exhibit \
        ↳four properties over time: 1) \
        constant mean (no trend), 2) constant variance (no heteroscedasticity), 3) \
        ↳constant autocorrelation structure, \
        4) no periodic component (no seasonality).")
print ("Above we discussed visually that the data seem to have to some extent a \
        ↳constant mean, a constant variance \
        a daily periodicity and some autocorrelation when the daily values are averaged. \
        ↳")

print("Now let's use some summary statistics and statistical tests to determine \
        ↳the stationarity.")
```

It is important to determine whether the Tair is stationary. This will determine which methods we will use to forecast it, and will tell whether we have to transform the data to get a cleaner signal (in non-stationary).

In order for a timeseries data to be stationary, the data must exhibit four properties over time: 1) constant mean (no trend), 2) constant variance (no heteroscedasticity), 3) constant autocorrelation structure, 4) no periodic component (no seasonality).

Above we discussed visually that the data seem to have to some extent a constant mean, a constant variance a daily periodicity and some autocorrelation when the daily values are averaged.

Now let's use some summary statistics and statistical tests to determine the stationarity.

```
[37]: Tair_hr_values = Tair_hr.values
chunks = np.array_split(Tair_hr_values[:, 0], 50)
```

```
chunks[0]
```

```
[37]: array([20.775      , 20.74166667, 20.925      , 21.00833333, 21.11666667,
          20.98333333, 21.09166667, 21.15      , 21.65      , 21.49166667,
          22.01666666, 22.05833333, 21.82500001, 21.68333333, 22.91666666,
          24.65833333, 23.35      , 20.81666667, 17.65833334, 14.53333334,
          13.80833333, 13.175      , 12.55833333, 11.95833333, 14.46666666,
          15.93333333, 16.175      , 16.41666666, 17.49166666, 17.725      ,
          17.675      , 17.85833333, 17.88333333, 18.13333333, 19.36666666,
          20.31666666, 20.95      , 22.19999999, 22.03333333, 21.24166667,
          20.85      , 20.50833333, 18.75000001, 15.56666667, 13.69166667,
          13.1      , 12.925      , 12.90833333, 14.85833332, 16.29166666,
          16.85833333, 16.91666666, 17.79166666, 18.51666666, 18.29166667,
          18.01666667, 18.10833333, 18.225      , 20.      , 21.3      ,
          21.48333333, 21.48333334, 21.05833333, 21.33333334, 21.15      ,
          20.58333333, 17.45000002, 13.70833333, 12.88333334, 12.39166667,
          12.18333333, 11.99166666, 14.48333332, 16.14166667, 16.83333333,
          17.15833333, 17.61666666, 18.73333333, 18.69166667, 18.51666667])
```

```
[38]: print("{} | {:.7} | {}".format("Chunk", "Mean", "Variance"))
      print("-" * 26)
      for i, chunk in enumerate(chunks, 1):
          print("{:5} | {:.6} | {:.6}".format(i, np.mean(chunk), np.var(chunk)))
```

Chunk	Mean	Variance
-------	------	----------

1	18.2396	10.4118
2	18.3018	9.2496
3	19.4288	15.9035
4	21.0478	17.9846
5	22.5196	20.286
6	22.1696	22.8943
7	22.6919	17.2257
8	23.3164	15.1268
9	22.6403	15.5943
10	23.0004	13.8433
11	23.2543	15.0368
12	22.6644	15.9071
13	22.6557	12.6901
14	23.0409	16.0373
15	22.746	16.9179
16	23.2919	16.5727
17	23.5159	15.0981
18	22.6602	15.4158
19	23.2593	13.5949
20	23.7118	16.1951



21		22.2999		13.6625
22		22.5501		10.5809
23		23.0509		13.9558
24		22.9685		19.048
25		22.1141		11.8366
26		22.8354		12.3359
27		22.467		14.4393
28		22.9278		13.1766
29		23.2932		13.7182
30		22.6157		17.0552
31		22.8194		13.3028
32		23.1356		12.6402
33		23.2129		10.8642
34		23.6071		8.64246
35		24.1937		9.29517
36		23.7237		9.49956
37		23.3552		9.74342
38		23.5098		8.22818
39		23.2488		11.2616
40		22.0418		10.9013
41		21.8893		8.7423
42		22.5912		10.6787
43		22.4973		11.6862
44		22.9328		13.7114
45		22.5078		9.08623
46		23.3632		9.26753
47		23.6667		12.643
48		24.4284		15.9175
49		23.7815		11.3102
50		24.1479		16.2739


```
[39]: print ("There is some small deviation from the mean, where the initial values
↪are a little smaller and the last ones \
a little higher, but not too much. There is some change in the variance, but it
↪is not too big. Let's do more tests \
before we drive conclusions.")
```

There is some small deviation from the mean, where the initial values are a little smaller and the last ones a little higher, but not too much. There is some change in the variance, but it is not too big. Let's do more tests before we drive conclusions.

```
[40]: Tair_hr_values = Tair_hr.values
Tair_hr_series = pd.Series(Tair_hr_values[:, 0])

Tair_hr_series.hist(bins=20)
```

```
[40]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9b3c0d7b50>
```



```
[41]: print ("This is not a Gaussian distribution. This indicates that the dataset is_\n\n↪non stationary. \n\nIn the histogram we can see the daily cicle of the data.")
```

This is not a Gaussian distribution. This indicates that the dataset is non stationary. In the histogram we can see the daily cycle of the data.

```
[42]: print("Let's run the augmented Dikey-Fuller statistical test to test the \n\nstationarity of our dataset. The null hypothesis is that the series is_\n\n↪nonstationary. We set a significance \n\nlevel to accept or reject the null to 0.05.")\n# lets transform the pandas dataframe to a numpy array\nTair_hr_values = Tair_hr.values\nfrom statsmodels.tsa.stattools import adfuller\n\nadf, pvalue, usedlag, nobs, critical_values, icbest = adfuller(Tair_hr_values[:\n\n↪, 0])\n\nprint("The adf value is ",adf)\nprint ("The more negative the value, the more confident we can be that the_\n\n↪series is stationary. But let's \n\nprint the p-value:",pvalue)\nprint ("The p-value is too small, so we reject the null that our data is_\n\n↪nonstationary")
```

```
print ("The Dickey-Fuller test indicates that the dataset is stationary, but we
    ↳do see a daily pattern in the data. \
This can happen because the test primarily focuses on detecting trends rather
    ↳than periodic patterns.")
```

Let's run the augmented Dickey-Fuller statistical test to test the stationarity of our dataset. The null hypothesis is that the series is nonstationary. We set a significance level to accept or reject the null to 0.05.

The adf value is -3.933187900689058

The more negative the value, the more confident we can be that the series is stationary. But let's print the p-value: 0.0018037339247767325

The p-value is too small, so we reject the null that our data is nonstationary

The Dickey-Fuller test indicates that the dataset is stationary, but we do see a daily pattern in the data. This can happen because the test primarily focuses on detecting trends rather than periodic patterns.

```
[43]: print ("Presence of seasonality: The ADF test is primarily designed to assess
    ↳the presence of a unit root and \
does not directly account for seasonality. If the data exhibits significant
    ↳seasonal patterns, the ADF test may \
fail to identify non-stationarity caused by seasonal components. Let's run an
    ↳ADF test that has a focus on seasonality")
```

Presence of seasonality: The ADF test is primarily designed to assess the presence of a unit root and does not directly account for seasonality. If the data exhibits significant seasonal patterns, the ADF test may fail to identify non-stationarity caused by seasonal components. Let's run an ADF test that has a focus on seasonality

```
[44]: from statsmodels.tsa.stattools import adfuller

# Assuming Tair_hr_values is your hourly dataset
Tair_hr_series = pd.Series(Tair_hr_values[:, 0])

# Perform the ADF test on the seasonal component
adf_result = adfuller(Tair_hr_series.diff(24).dropna())

# Extract the p-value
p_value = adf_result[1]

# Print the test result and explanation
print("ADF Test Result:")
if p_value < 0.05:
    print("The dataset is not stationary, indicating the presence of a daily
        ↳cycle.")
else:
```

```
print("The dataset is stationary.")
```

ADF Test Result:

The dataset is not stationary, indicating the presence of a daily cycle.

```
[45]: from statsmodels.tsa.stattools import kpss

# Perform KPSS test
result = kpss(Tair_hr_values[:, 0])

# Extract the p-value from the test result
p_value = result[1]

# Compare the p-value with a significance level (e.g., 0.05) to determine
↳stationarity
if p_value < 0.05:
    print("The time series is not stationary.")
else:
    print("The time series is stationary.")
print (p_value)
```

The time series is not stationary.

0.01

```
[46]: comment = "**4. Univariate Temperature forecasting using smoothing**"
font_size = "20px"

Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</
↳span>")
```

[46]: 4. Univariate Temperature forecasting using smoothing

```
[47]: print ("We will forecast the last 24 hours of the indoors temperature using:")
print ("1. Simple Average Smoothing")
print ("2. Triple Exponential Smoothing")

print ("We won't use the single exponential smoothing because it is mainly
↳useful for short-term \
forecasting and when the data has a gradual or linear trend that needs to be
↳captured. \
Likewise, we skip the double exponential smoothing because it refers to
↳timeseries \
with a trend. We have a look at the triplpe exponential smoothing that refers
↳to \
timeseries with trend and seasonality. There is no significant trend in our
↳data though.")
```

We will forecast the last 24 hours of the indoors temperature using:

1. Simple Average Smoothing
2. Triple Exponential Smoothing

We won't use the single exponential smoothing because it is mainly useful for short-term forecasting and when the data has a gradual or linear trend that needs to be captured. Likewise, we skip the double exponential smoothing because it refers to timeseries with a trend. We have a look at the triplpe exponential smoothing that refers to timeseries with trend and seasonality. There is no significant trend in our data though.

```
[48]: train = Tair_hr_values[:-24]
      test = Tair_hr_values[-24:]
      train[3900:-24],test
```

```
[48]: (array([[31.16666667],
              [31.86666667],
              [29.38333333],
              [28.18333333],
              [29.00833333],
              [27.44166667],
              [25.00833334],
              [22.73333333],
              [21.7       ],
              [21.01666667],
              [20.34583333],
              [20.66666666],
              [20.74166667],
              [20.48333333],
              [20.70833333],
              [20.74166667],
              [20.275     ],
              [19.88333333],
              [20.58333333],
              [22.55833333],
              [24.4       ],
              [26.49166666],
              [27.73333333],
              [29.75      ],
              [30.6       ],
              [30.70833333],
              [29.89166666],
              [28.61666667],
              [28.65000002],
              [27.93333334],
              [25.39166667],
              [22.75000001],
              [21.39166666],
              [21.10833333],
```

```

[21.06666666],
[20.86666667]), array([[20.31666667],
[20.23333333],
[20.325      ],
[20.43333333],
[20.45      ],
[20.075     ],
[20.85      ],
[22.1       ],
[24.54166666],
[26.75      ],
[28.275     ],
[29.125     ],
[30.66666667],
[31.01666667],
[30.875     ],
[29.82500001],
[28.15000001],
[26.40833333],
[24.55833334],
[22.91666667],
[21.44166667],
[19.8       ],
[18.77916666],
[18.35833334]]))

```

```

[49]: def mse(observations, estimates):
    '''
    INPUT:
        observations - numpy array of values indicating observed values
        estimates - numpy array of values indicating an estimate of values
    OUTPUT:
        Mean Square Error value
    '''
    # check arg types
    assert type(observations) == type(np.array([])), "'observations' must be a_
↳numpy array"
    assert type(estimates) == type(np.array([])), "'estimates' must be a numpy_
↳array"
    # check length of arrays equal
    assert len(observations) == len(estimates), "Arrays must be of equal length"

    # calculations
    difference = observations - estimates
    sq_diff = difference ** 2
    mse = sum(sq_diff)

```

```
return mse
```

```
[50]: # find mean of series
trend_seasonal_avg = np.mean(Tair_hr_values)

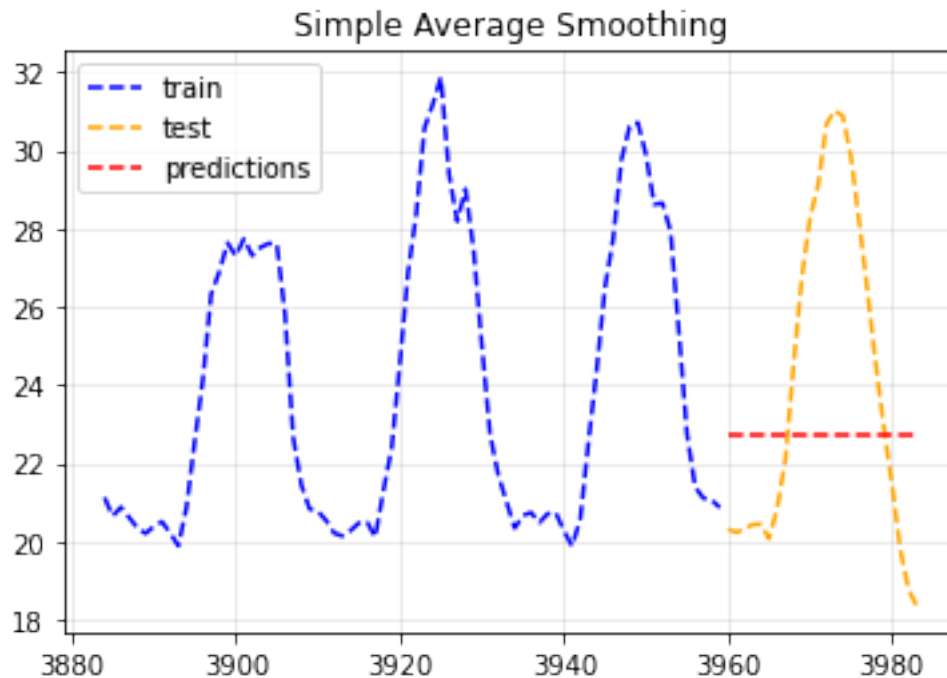
# create array of mean value equal to length of time array
simple_avg_preds = np.full(shape=len(test), fill_value=trend_seasonal_avg,
    dtype='float')

# mse
simple_mse = mse(test, simple_avg_preds)

# results
print("Predictions: ", simple_avg_preds)
print("MSE: ", simple_mse)
```

```
Predictions: [22.71656505 22.71656505 22.71656505 22.71656505 22.71656505
22.71656505
22.71656505 22.71656505 22.71656505 22.71656505 22.71656505 22.71656505
22.71656505 22.71656505 22.71656505 22.71656505 22.71656505 22.71656505]
MSE: [470.80015169 470.80015169 470.80015169 470.80015169 470.80015169
470.80015169 470.80015169 470.80015169 470.80015169 470.80015169
470.80015169 470.80015169 470.80015169 470.80015169 470.80015169
470.80015169 470.80015169 470.80015169 470.80015169]
```

```
[51]: time = np.arange(3984)
train_subset = train[-100:].flatten()
plt.plot(time[-100:-24], train_subset[:24], 'b--', label="train")
plt.plot(time[-24:], test, color='orange', linestyle="--", label="test")
plt.plot(time[-24:], simple_avg_preds, 'r--', label="predictions")
plt.legend(loc='upper left')
plt.title("Simple Average Smoothing")
plt.grid(alpha=0.3)
```



[ ]:

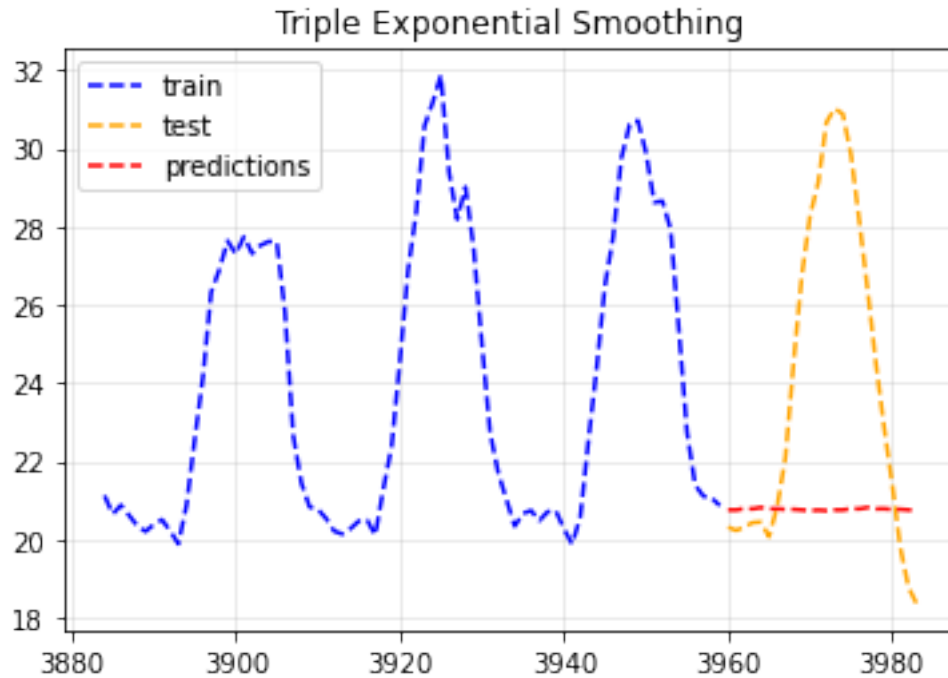
[52]: `from statsmodels.tsa.api import ExponentialSmoothing`

```
triple = ExponentialSmoothing(train,
                              trend="additive",
                              seasonal="additive",
                              seasonal_periods=13).fit(optimized=True)
triple_preds = triple.forecast(len(test))
triple_mse = mse(test, triple_preds)
print("Predictions: ", triple_preds)
print("MSE: ", triple_mse)
```

```
Predictions: [20.75802878 20.75785392 20.79466218 20.7856051 20.82008972
20.78557375
20.78921815 20.77163693 20.77400625 20.75883463 20.74187863 20.75390028
20.74171441 20.75807602 20.75790117 20.79470943 20.78565234 20.82013697
20.78562099 20.7892654 20.77168417 20.7740535 20.75888188 20.74192588]
MSE: [684.57715711 684.60446223 678.88882858 680.28919212 674.9783876
680.29404666 679.73008658 682.45660366 682.08830084 684.45133315
687.10536361 685.22225721 687.13113646 684.56977963 684.59708436
678.88153418 680.28187718 674.97115086 680.28673164 679.72277983
682.44925704 682.08095959 684.4439575 687.09794951]
```



```
[53]: time = np.arange(3984)
train_subset = train[-100:].flatten()
plt.plot(time[-100:-24], train_subset[:24], 'b--', label="train")
plt.plot(time[-24:], test, color='orange', linestyle="--", label="test")
plt.plot(time[-24:], triple_preds, 'r--', label="predictions")
plt.legend(loc='upper left')
plt.title("Triple Exponential Smoothing")
plt.grid(alpha=0.3);
```



```
[54]: print ("Indeed it does not do a very good job at predicting temperature values.
↪")
```

Indeed it does not do a very good job at predicting temperature values.

```
[55]: print("COMPARISON")
data_dict = {'MSE': [np.mean(simple_mse), np.mean(triple_mse)]}
df = pd.DataFrame(data_dict, index=['simple', 'triple'])
print(df)
```

COMPARISON

```

              MSE
simple 470.80015
triple 682.13334
```

```
[56]: comment = "**5. Univariate Temperature forecasting using SARIMA Model**"
font_size = "20px"

Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</span>")
```

[56]: 5. Univariate Temperature forecasting using SARIMA Model

```
[57]: print ("We are going to use a SARIMA model to predict indoors temperature. The
↳SARIMA stands \
for Seasonal Autoregressive Integrated Moving Average Model.")
```

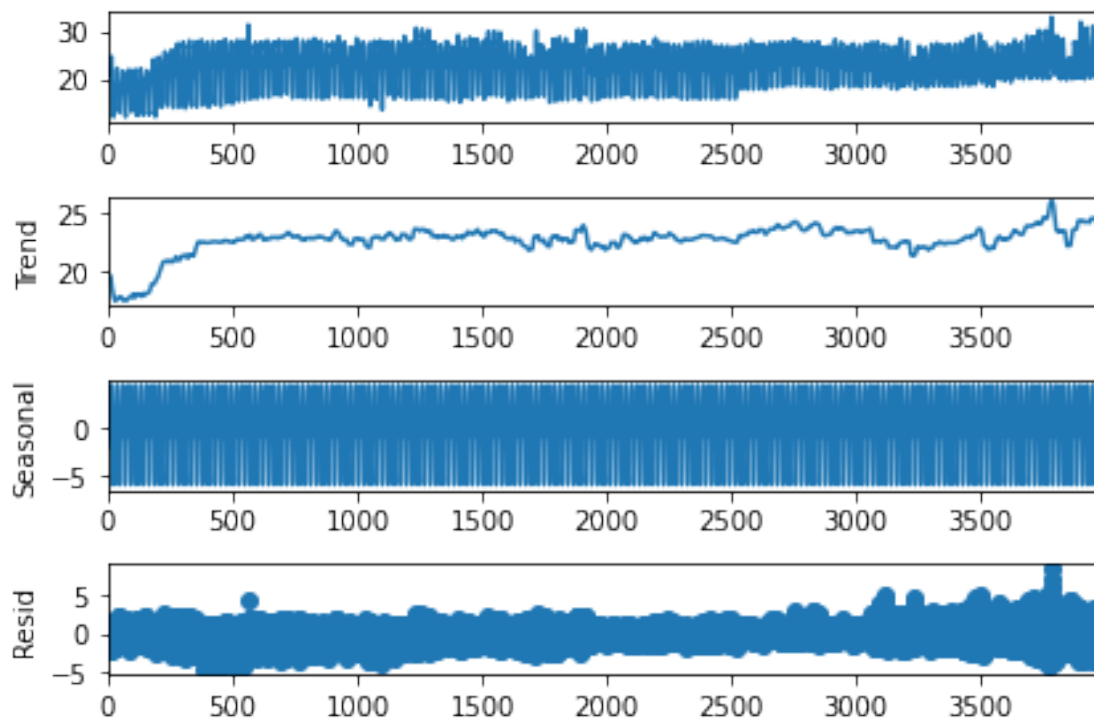
We are going to use a SARIMA model to predict indoors temperature. The SARIMA stands for Seasonal Autoregressive Integrated Moving Average Model.

```
[58]: """
order(p,d,q):
p is number of AR terms
d is number of times that we would difference our data
q is number of MA terms
When we work with SARIMA models 'S' refers to 'seasonal' and we have the
↳additional
standard inputs:


seasonal order(p,d,q):
p is number of AR terms in regards to seasonal lag
d is number of times that we would difference our seasonal lag (as seen above)
q is number of MA terms in regards to seasonal lag
s is number of periods in a season
"""
```

```
[58]: "\norder(p,d,q):\np is number of AR terms\nd is number of times that we would
difference our data\nq is number of MA terms\nWhen we work with SARIMA models
'S' refers to 'seasonal' and we have the additional \nstandard
inputs:\n\nseasonal order(p,d,q):\np is number of AR terms in regards to
seasonal lag\nd is number of times that we would difference our seasonal lag (as
seen above)\nq is number of MA terms in regards to seasonal lag\ns is number of
periods in a season\n"
```

```
[59]: from statsmodels.tsa.seasonal import seasonal_decompose
seasonal_decompose(Tair_hr[:],period=24).plot();
```





```
[60]: from statsmodels.tsa.seasonal import seasonal_decompose
seasonal_decompose(Tair_hr[3700:],period=24).plot();
```



```
[61]: import statsmodels.api as sm

sm.tsa.graphics.plot_acf(Tair_hr[3000:],zero=False)
sm.tsa.graphics.plot_pacf(Tair_hr[3000:],zero = False);
```





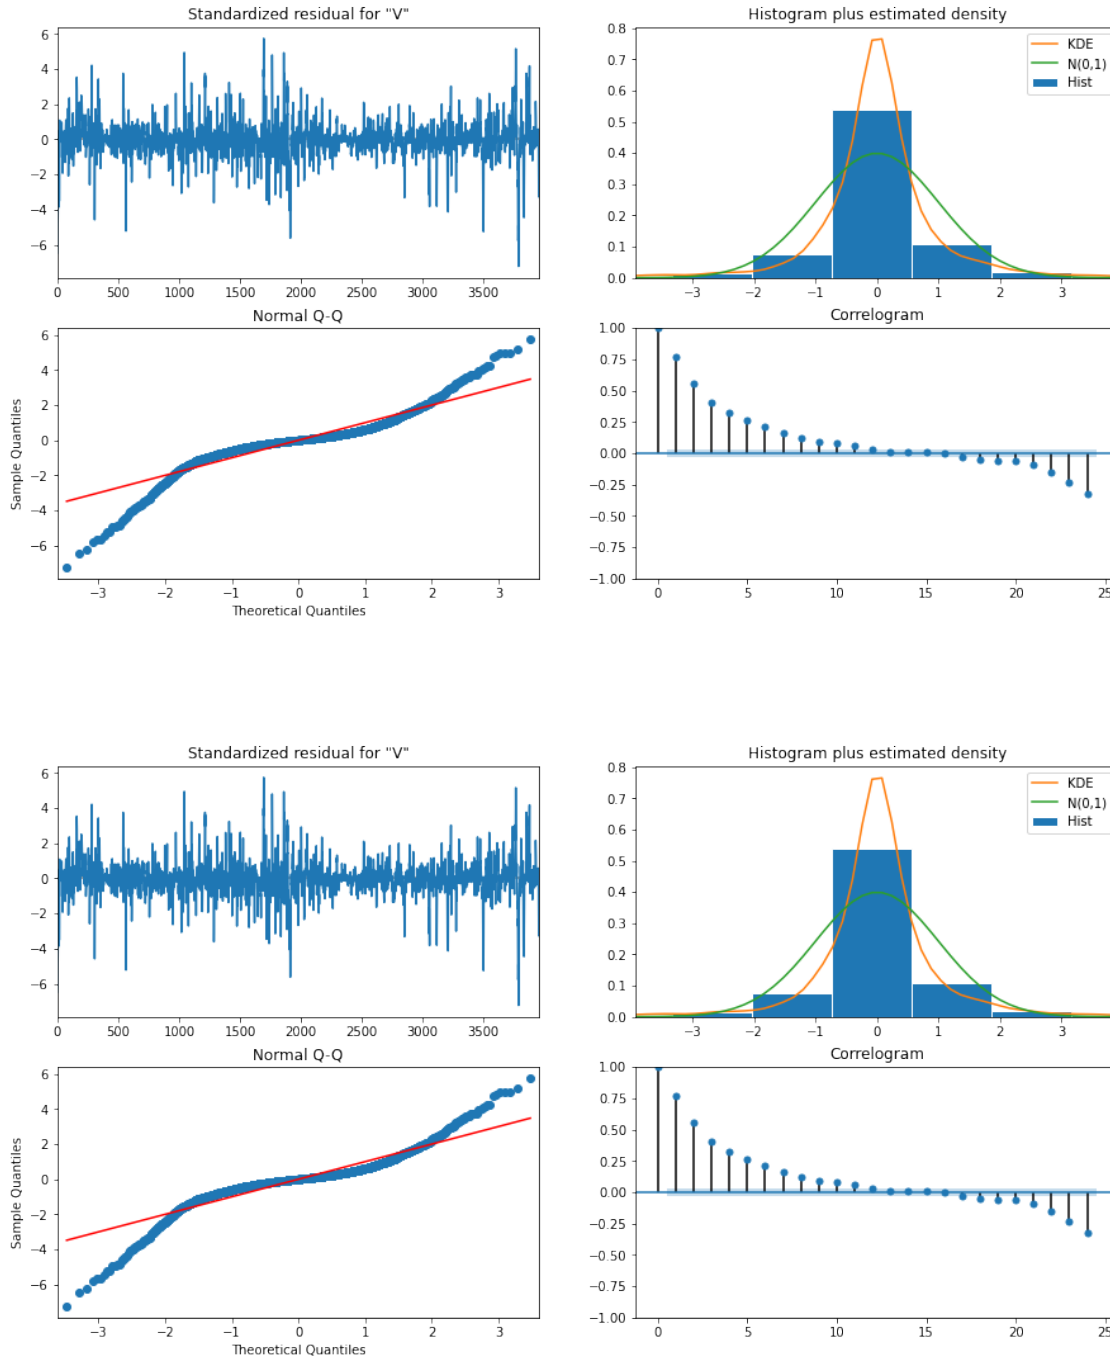
```
[62]: print ("Now let's start with the most basic model, with no auto correlation and
↳we just \
difference the values and predict moving forward.")

sar = sm.tsa.statespace.SARIMAX(Tair_hr.iloc[:, 0],
                                order=(0,0,0),
                                seasonal_order=(0,1,0,24),
                                trend='c').fit()
```

Now let's start with the most basic model, with no auto correlation and we just difference the values and predict moving forward.

```
[63]: sar.plot_diagnostics(figsize = (15,8),lags=24)
```

```
[63]:
```



```
[64]: print ("The pattern in the residuals suggests that the model can be further_
↪improved. \
The histogram suggests that the standardized residuals follow to some_
↪approximation the \
normal distribution. The plot compares the quantiles of the observed data_
↪against the \
```

```
quantiles of a theoretical standard normal distribution. We see some skewness,
→and some \
outliers. The correlogram suggests significant autocorrelation because the
→autocorrelation \
falls really outside the confidence intervals.")
```

The pattern in the residuals suggests that the model can be further improved. The histogram suggests that the standardized residuals follow to some approximation the normal distribution. The plot compares the quantiles of the observed data against the quantiles of a theoretical standard normal distribution. We see some skewness and some outliers. The correlogram suggests significant autocorrelation because the autocorrelation falls really outside the confidence intervals.

```
[65]: Tair_hr_values = Tair_hr.values
      Tair_hr.iloc[2000:, 0]
```

```
[65]: 2000    23.24167
      2001    24.37500
      2002    24.10833
      2003    24.58333
      2004    25.39167
      ...
      3979    22.91667
      3980    21.44167
      3981    19.80000
      3982    18.77917
      3983    18.35833
      Name: Value, Length: 1984, dtype: float64
```

```
[66]: Tair_hr.index
```

```
[66]: RangeIndex(start=0, stop=3984, step=1)
```

```
[67]: print ("for the next steps we need to set a datetime to the data")
      start_date = '2013-01-01'
      end_date = '2023-11-28'
      step = pd.DateOffset(days=1)

      Tair_hr_index = pd.date_range(start=start_date, end=end_date, freq=step)

      # Print the first few values
      print(Tair_hr_index[:5])
```

```
for the next steps we need to set a datetime to the data
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
              '2013-01-05'],
```

```
dtype='datetime64[ns]', freq='<DateOffset: days=1>')
```

```
[69]: Tair_hr.index = Tair_hr_index
Tair_hr.index
```

```
[69]: DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
                    '2013-01-05', '2013-01-06', '2013-01-07', '2013-01-08',
                    '2013-01-09', '2013-01-10',
                    ...,
                    '2023-11-19', '2023-11-20', '2023-11-21', '2023-11-22',
                    '2023-11-23', '2023-11-24', '2023-11-25', '2023-11-26',
                    '2023-11-27', '2023-11-28'],
                    dtype='datetime64[ns]', length=3984, freq='<DateOffset: days=1>')
```

```
[70]: df_train = Tair_hr.iloc[3500:-24, 0]
df_test = Tair_hr.iloc[-24:, 0]
df_train.tail()
```

```
[70]: 2023-10-31    23.13333
2023-11-01    21.36667
2023-11-02    20.79167
2023-11-03    21.26667
2023-11-04    20.74167
Freq: <DateOffset: days=1>, Name: Value, dtype: float64
```

```
[71]: df_test.head()
```

```
[71]: 2023-11-05    20.31667
2023-11-06    20.23333
2023-11-07    20.32500
2023-11-08    20.43333
2023-11-09    20.45000
Freq: <DateOffset: days=1>, Name: Value, dtype: float64
```

```
[72]: print ("Let's run auto ARIMA that will give us the best fitting terms that \
        ↳result in the \
smallest AIC (Akaike Information Criterion, a metric used to compare the \
        ↳performance \
of different models).")
print("I have commented out the lines below because it takes minutes to run \
the script and we don't want to wait during the interview :)")
import pmdarima as pm
# Seasonal - fit stepwise auto-ARIMA
SARIMA_model = pm.auto_arima(df_train, start_p=0, start_q=0,
                             test='adf',
                             max_p=3, max_q=3,
                             m=24, #12 is the frequency of the cycle
```



```

        start_P=0,
        seasonal=True, #set to seasonal
        d=None,
        D=1, #order of the seasonal differencing
        trace=False,
        error_action='ignore',
        suppress_warnings=True,
        stepwise=True)
print("Best Parameters (p, d, q, P, D, Q, s):", SARIMA_model.order,
      ↪SARIMA_model.seasonal_order)
print(SARIMA_model.aic())

```

Let's run auto ARIMA that will give us the best fitting terms that result in the smallest AIC (Akaike Information Criterion, a metric used to compare the performance of different models).

I have commented out the lines below because it takes minutes to run the script and we don't want to wait during the interview :)

```

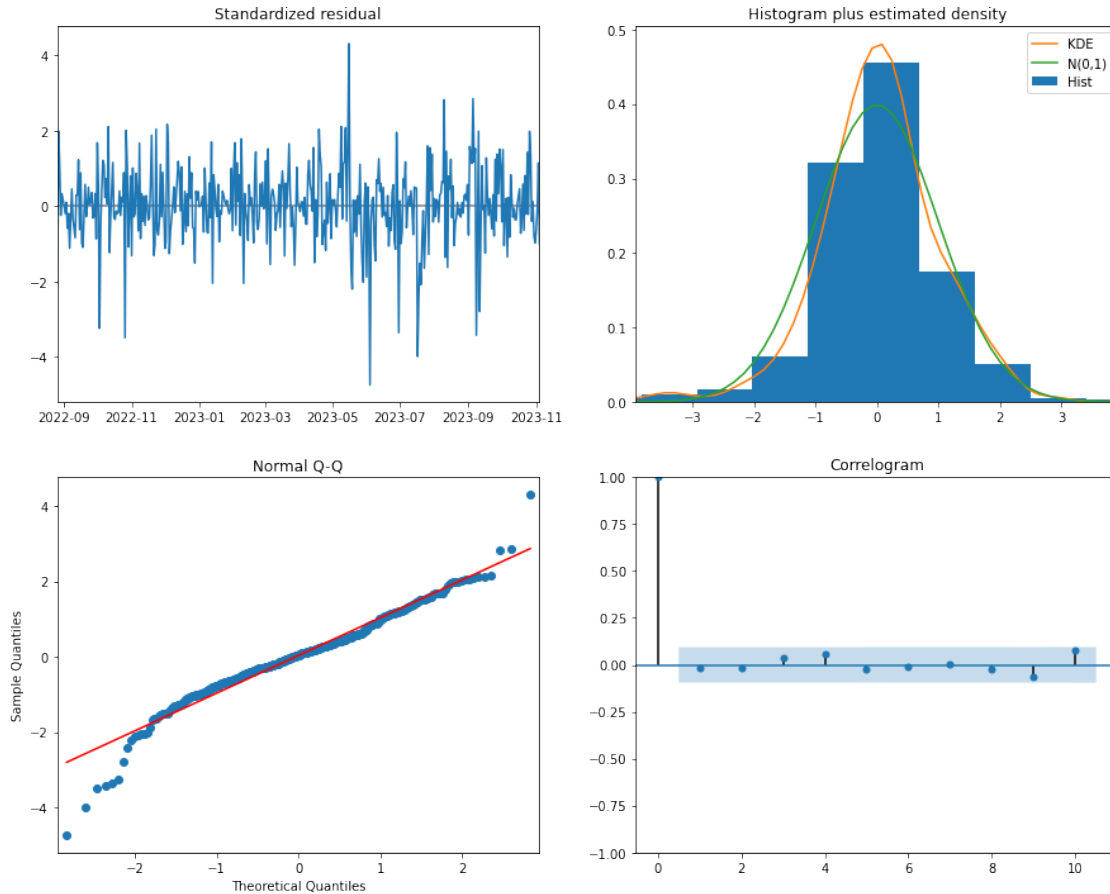
Best Parameters (p, d, q, P, D, Q, s): (2, 0, 1) (2, 1, 0, 24)
735.4039780741509

```

```

[75]: SARIMA_model.plot_diagnostics(figsize=(15,12))
      plt.show()

```



```
[76]: print ("Here we see more like a random noise in the residuals, a fairly normal_
↪distribution \
and almost no autocorrelation.")
```

Here we see more like a random noise in the residuals, a fairly normal distribution and almost no autocorrelation.

```
[77]: def forecast(SARIMA_model, periods=24):
    # Forecast
    n_periods = periods
    fitted, confint = SARIMA_model.predict(n_periods=n_periods,
↪return_conf_int=True)
    index_of_fc = pd.date_range(df_train.index[-1] + pd.DateOffset(days=1),
↪periods = n_periods, freq='D')

    forecast_values_auto = fitted
    #confidence_intervals_auto = conf_int_auto
```

```


# make series for plotting purpose
fitted_series = pd.Series(fitted, index=index_of_fc)
lower_series = pd.Series(confint[:, 0], index=index_of_fc)
upper_series = pd.Series(confint[:, 1], index=index_of_fc)
forecast_auto = fitted_series.values
lower_percentile = np.percentile(confint, 10, axis=1)
upper_percentile = np.percentile(confint, 90, axis=1)

# Plot
plt.figure(figsize=(15,7))
plt.plot(df_train[-50:], color='#1f76b4', label='Training')
plt.plot(fitted_series, color='red', label='Predicted')
plt.plot(df_test, color='black', label='Actual')
plt.fill_between(lower_series.index, lower_percentile, upper_percentile,
color='k', alpha=.15, label='90% Confidence Interval')

plt.title("SARIMA - Forecast of indoors Temperature")
plt.legend(loc='upper left')
plt.show()

forecast(SARIMA_model)

```



```
[78]: SARIMA_model.summary()
```

```
[78]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

SARIMAX Results

=====

```

=====
Dep. Variable:          y    No. Observations:
460
Model:                SARIMAX(2, 0, 1)x(2, 1, [], 24)    Log Likelihood
-361.702
Date:                  Sun, 09 Jul 2023    AIC
735.404
Time:                  14:50:39    BIC
759.870
Sample:                08-02-2022    HQIC
745.059

                                - 11-04-2023
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.6469      0.134      4.826      0.000      0.384      0.910
ar.L2          0.1994      0.123      1.616      0.106     -0.042      0.441
ma.L1          0.5068      0.119      4.255      0.000      0.273      0.740
ar.S.L24       -0.5121      0.036     -14.065      0.000     -0.583     -0.441
ar.S.L48       -0.2994      0.046      -6.535      0.000     -0.389     -0.210
sigma2         0.3004      0.014     21.221      0.000      0.273      0.328
=====
===
Ljung-Box (L1) (Q):          0.16    Jarque-Bera (JB):
158.97
Prob(Q):          0.68    Prob(JB):
0.00
Heteroskedasticity (H):      1.34    Skew:
-0.43
Prob(H) (two-sided):        0.08    Kurtosis:
5.83
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""

```

```

[79]: # Predicted values and actual values
forecast_auto, conf_int_auto = SARIMA_model.predict(n_periods=24,
                                                    return_conf_int=True)

predicted_values_auto = forecast_auto.values
actual_values = df_test.values.flatten()

# Mean Absolute Error (MAE)

```

```

maeSA = np.mean(np.abs(predicted_values_auto - actual_values))
print("MAE:", maeSA)

# Root Mean Squared Error (RMSE)
mseSA = np.mean((predicted_values_auto - actual_values) ** 2)
rmseSA = np.sqrt(mseSA)
print("MSE:", rmseSA)
print("RMSE:", rmseSA)

# Mean Absolute Percentage Error (MAPE)
mapeSA = np.mean(np.abs((predicted_values_auto - actual_values) /
    ↪actual_values)) * 100
print("MAPE:", mapeSA)

```

MAE: 0.5940614094719291  
MSE: 0.9558774709361916  
RMSE: 0.9558774709361916  
MAPE: 2.6084199193111255

```

[80]: comment = "***6. Multivariate Temperature forecasting using SARIMAX Model**"
font_size = "20px"

Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</
    ↪span>")

```

[80]: 6. Multivariate Temperature forecasting using SARIMAX Model

```

[81]: print("SARIMAX(Seasonal Auto-Regressive Integrated Moving Average with
    ↪eXogenous factors) \
is an updated version of the ARIMA model. ARIMA includes an autoregressive
    ↪integrated moving \
average, while SARIMAX includes seasonal effects and eXogenous factors with the
    ↪\
autoregressive and moving average component in the model.")

```

SARIMAX(Seasonal Auto-Regressive Integrated Moving Average with eXogenous factors) is an updated version of the ARIMA model. ARIMA includes an autoregressive integrated moving average, while SARIMAX includes seasonal effects and eXogenous factors with the autoregressive and moving average component in the model.

```

[82]: print ("we have lots of variables that could influence the indoors temperature.
    ↪Some of \
them are clearly not important (eg drain pH and drain EC). For the rest we can
    ↪to a feature \
importance analysis, to see which among the variables explain most of the
    ↪variance")

```

we have lots of variables that could influence the indoors temperature. Some of them are clearly not important (eg drain pH and drain EC). For the rest we can to a feature importance analysis, to see which among the variables explain most of the variance

```
[83]: df_outd.head()
```

```
[83]:
```

	time	AbsHumOut	Iglob	PARout	Pyrgeo	RadSum	Rain	\
0	43,815.00000	6.22095	0.00000	0.00000	-72.00000	215.00000	0.00000	
1	43,815.00347	6.22095	0.00000	0.00000	-73.00000	0.00000	0.00000	
2	43,815.00694	6.20556	0.00000	0.00000	-76.00000	0.00000	0.00000	
3	43,815.01042	6.19017	0.00000	0.00000	-77.00000	0.00000	0.00000	
4	43,815.01389	6.16262	0.00000	0.00000	-75.00000	0.00000	0.00000	

	Rhout	Tout	Winddir	Windsp
0	80.60000	6.90000	32.00000	4.70000
1	80.60000	6.90000	32.00000	4.70000
2	80.40000	6.90000	32.00000	4.70000
3	80.20000	6.90000	32.00000	4.70000
4	80.90000	6.70000	32.00000	4.70000

```
[84]: df_ind.head()
```

```
[84]:
```

	time	AssimLight	BlackScr	CO2air	Cum_irr	EC_drain_PC	\
0	43,815.00000	0.00000	35.00000	472.00000	0.35480	3.98000	
1	43,815.00347	0.00000	85.00000	501.00000	0.53220	3.99000	
2	43,815.00694	0.00000	96.00000	489.00000	0.53220	3.99000	
3	43,815.01042	0.00000	96.00000	497.00000	0.53220	4.00000	
4	43,815.01389	0.00000	96.00000	477.00000	0.53220	4.00000	

	EnScr	HumDef	PipeGrow	PipeLow	Rhair	Tair	Tot_PAR	\
0	100.00000	7.17000	0.00000	45.20000	60.60000	20.90000	0.00000	
1	100.00000	6.94000	0.00000	43.60000	61.40000	20.70000	0.00000	
2	100.00000	7.24000	0.00000	42.30000	60.90000	21.20000	0.00000	
3	100.00000	6.74000	0.00000	41.20000	62.50000	20.70000	0.00000	
4	100.00000	6.71000	0.00000	41.20000	62.70000	20.70000	0.00000	

	Tot_PAR_Lamps	VentLee	Ventwind	assim_sp	assim_vip	co2_dos	co2_sp	\
0	0.00000	0.80000	0.00000	NaN	0.00000	NaN	NaN	
1	0.00000	0.80000	0.00000	NaN	0.00000	0.00034	NaN	
2	0.00000	0.80000	0.00000	NaN	0.00000	0.00036	NaN	
3	0.00000	0.80000	0.00000	NaN	0.00000	0.00201	NaN	
4	0.00000	0.80000	0.00000	NaN	0.00000	0.00155	NaN	

	co2_vip	dx_sp	dx_vip	int_blue_sp	int_blue_vip	int_farred_sp	\
0	400.00000	NaN	2.20000	NaN	NaN	NaN	
1	400.00000	NaN	2.20000	0.00000	NaN	0.00000	

2	400.00000	NaN	2.20000	0.00000	NaN	0.00000
3	400.00000	NaN	2.20000	0.00000	NaN	0.00000
4	400.00000	NaN	2.20000	0.00000	NaN	0.00000

	int_farred_vip	int_red_sp	int_red_vip	int_white_sp	int_white_vip	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	0.00000	NaN	0.00000	NaN	
2	NaN	0.00000	NaN	0.00000	NaN	
3	NaN	0.00000	NaN	0.00000	NaN	
4	NaN	0.00000	NaN	0.00000	NaN	

	pH_drain_PC	scr_blk_sp	scr_blk_vip	scr_enrg_sp	scr_enrg_vip	\
0	6.27000	NaN	96.00000	NaN	100.00000	
1	6.28000	NaN	96.00000	NaN	100.00000	
2	6.28000	NaN	96.00000	NaN	100.00000	
3	6.27000	NaN	96.00000	NaN	100.00000	
4	6.28000	NaN	96.00000	NaN	100.00000	

	t_grow_min_sp	t_grow_min_vip	t_heat_sp	t_heat_vip	t_rail_min_sp	\
0	NaN	10.00000	NaN	21.00000	NaN	
1	NaN	10.00000	NaN	21.00000	NaN	
2	NaN	10.00000	NaN	21.00000	NaN	
3	NaN	10.00000	NaN	21.00000	NaN	
4	NaN	10.00000	NaN	21.00000	NaN	

	t_rail_min_vip	t_vent_sp	t_ventlee_vip	t_ventwind_vip	water_sup	\
0	0.00000	NaN	25.00000	26.00000	4.00000	
1	0.00000	NaN	25.00000	26.00000	6.00000	
2	0.00000	NaN	25.00000	26.00000	6.00000	
3	0.00000	NaN	25.00000	26.00000	6.00000	
4	0.00000	NaN	25.00000	26.00000	6.00000	

	water_sup_intervals_sp_min	water_sup_intervals_vip_min	window_pos_lee_sp	\
0	NaN	1,000.00000	NaN	
1	NaN	1,000.00000	NaN	
2	NaN	1,000.00000	NaN	
3	NaN	1,000.00000	NaN	
4	NaN	1,000.00000	NaN	

	window_pos_lee_vip
0	1.20000
1	1.20000
2	1.20000
3	1.20000
4	1.20000

```
[85]: print ("we have to combine the dataframes:")
df_combined = pd.concat([df_outd, df_ind], axis=1)
df_combined.head()
```

we have to combine the dataframes:

```
[85]:
```

	time	AbsHumOut	Iglob	PARout	Pyrgeo	RadSum	Rain	\
0	43,815.00000	6.22095	0.00000	0.00000	-72.00000	215.00000	0.00000	
1	43,815.00347	6.22095	0.00000	0.00000	-73.00000	0.00000	0.00000	
2	43,815.00694	6.20556	0.00000	0.00000	-76.00000	0.00000	0.00000	
3	43,815.01042	6.19017	0.00000	0.00000	-77.00000	0.00000	0.00000	
4	43,815.01389	6.16262	0.00000	0.00000	-75.00000	0.00000	0.00000	

	Rhout	Tout	Winddir	Windsp	time	AssimLight	BlackScr	\
0	80.60000	6.90000	32.00000	4.70000	43,815.00000	0.00000	35.00000	
1	80.60000	6.90000	32.00000	4.70000	43,815.00347	0.00000	85.00000	
2	80.40000	6.90000	32.00000	4.70000	43,815.00694	0.00000	96.00000	
3	80.20000	6.90000	32.00000	4.70000	43,815.01042	0.00000	96.00000	
4	80.90000	6.70000	32.00000	4.70000	43,815.01389	0.00000	96.00000	

	CO2air	Cum_irr	EC_drain_PC	EnScr	HumDef	PipeGrow	PipeLow	\
0	472.00000	0.35480	3.98000	100.00000	7.17000	0.00000	45.20000	
1	501.00000	0.53220	3.99000	100.00000	6.94000	0.00000	43.60000	
2	489.00000	0.53220	3.99000	100.00000	7.24000	0.00000	42.30000	
3	497.00000	0.53220	4.00000	100.00000	6.74000	0.00000	41.20000	
4	477.00000	0.53220	4.00000	100.00000	6.71000	0.00000	41.20000	

	Rhair	Tair	Tot_PAR	Tot_PAR_Lamps	VentLee	Ventwind	assim_sp	\
0	60.60000	20.90000	0.00000	0.00000	0.80000	0.00000	NaN	
1	61.40000	20.70000	0.00000	0.00000	0.80000	0.00000	NaN	
2	60.90000	21.20000	0.00000	0.00000	0.80000	0.00000	NaN	
3	62.50000	20.70000	0.00000	0.00000	0.80000	0.00000	NaN	
4	62.70000	20.70000	0.00000	0.00000	0.80000	0.00000	NaN	

	assim_vip	co2_dos	co2_sp	co2_vip	dx_sp	dx_vip	int_blue_sp	\
0	0.00000	NaN	NaN	400.00000	NaN	2.20000	NaN	
1	0.00000	0.00034	NaN	400.00000	NaN	2.20000	0.00000	
2	0.00000	0.00036	NaN	400.00000	NaN	2.20000	0.00000	
3	0.00000	0.00201	NaN	400.00000	NaN	2.20000	0.00000	
4	0.00000	0.00155	NaN	400.00000	NaN	2.20000	0.00000	

	int_blue_vip	int_farred_sp	int_farred_vip	int_red_sp	int_red_vip	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	0.00000	NaN	0.00000	NaN	
2	NaN	0.00000	NaN	0.00000	NaN	
3	NaN	0.00000	NaN	0.00000	NaN	
4	NaN	0.00000	NaN	0.00000	NaN	



	int_white_sp	int_white_vip	pH_drain_PC	scr_blkck_sp	scr_blkck_vip	\
0	NaN	NaN	6.27000	NaN	96.00000	
1	0.00000	NaN	6.28000	NaN	96.00000	
2	0.00000	NaN	6.28000	NaN	96.00000	
3	0.00000	NaN	6.27000	NaN	96.00000	
4	0.00000	NaN	6.28000	NaN	96.00000	

	scr_enrg_sp	scr_enrg_vip	t_grow_min_sp	t_grow_min_vip	t_heat_sp	\
0	NaN	100.00000	NaN	10.00000	NaN	
1	NaN	100.00000	NaN	10.00000	NaN	
2	NaN	100.00000	NaN	10.00000	NaN	
3	NaN	100.00000	NaN	10.00000	NaN	
4	NaN	100.00000	NaN	10.00000	NaN	

	t_heat_vip	t_rail_min_sp	t_rail_min_vip	t_vent_sp	t_ventlee_vip	\
0	21.00000	NaN	0.00000	NaN	25.00000	
1	21.00000	NaN	0.00000	NaN	25.00000	
2	21.00000	NaN	0.00000	NaN	25.00000	
3	21.00000	NaN	0.00000	NaN	25.00000	
4	21.00000	NaN	0.00000	NaN	25.00000	

	t_ventwind_vip	water_sup	water_sup_intervals_sp_min	\
0	26.00000	4.00000	NaN	
1	26.00000	6.00000	NaN	
2	26.00000	6.00000	NaN	
3	26.00000	6.00000	NaN	
4	26.00000	6.00000	NaN	

	water_sup_intervals_vip_min	window_pos_lee_sp	window_pos_lee_vip
0	1,000.00000	NaN	1.20000
1	1,000.00000	NaN	1.20000
2	1,000.00000	NaN	1.20000
3	1,000.00000	NaN	1.20000
4	1,000.00000	NaN	1.20000

```
[86]: print ("drop variables with many NaN values and then drop the Nan lines. Then,
        ↳take the \
        hourly values")
df_combined_filtered = df_combined.dropna(axis=1, thresh=df_combined.
        ↳shape[0]-150)
df_combined_filtered = df_combined_filtered.dropna()
num_groups = len(df_combined_filtered) // 12

# Take the average every 12 rows
df_combined_hr = df_combined_filtered.iloc[:num_groups*12].
        ↳groupby(df_combined_filtered.index // 12).mean()
```

```
df_combined_hr
```

drop variables with many NaN values and then drop the Nan lines. Then take the hourly values

```
[86]:
```

	time	AbsHumOut	Iglob	PARout	Pyrgeo	RadSum	\
0	43,815.02083	6.22694	0.00000	0.00000	-72.09091	0.00000	
1	43,815.06076	6.23459	0.00000	0.00000	-75.33333	0.00000	
2	43,815.10243	6.28361	0.00000	0.00000	-59.16667	0.00000	
3	43,815.14410	6.45165	0.00000	0.00000	-52.75000	0.00000	
4	43,815.18576	6.56769	0.00000	0.00000	-53.33333	0.00000	
...	...	...	...	...	...	...	
3979	43,980.81076	7.10976	276.75000	526.83333	-104.33333	2,896.75000	
3980	43,980.85243	7.47256	121.66667	235.33333	-99.83333	2,966.66667	
3981	43,980.89410	9.34072	20.41667	44.91667	-89.33333	2,989.83333	
3982	43,980.93576	9.40363	0.00000	1.00000	-85.25000	2,992.00000	
3983	43,980.97743	9.38771	0.00000	0.33333	-84.75000	2,992.00000	

	Rain	Rhout	Tout	Winddir	Windsp	time	AssimLight	\
0	0.00000	81.30909	6.78182	24.72727	4.06364	43,815.02083	0.00000	
1	0.00000	84.04167	6.30000	16.00000	2.90833	43,815.06076	0.00000	
2	0.00000	85.59167	6.14167	16.00000	3.20000	43,815.10243	0.00000	
3	0.00000	88.22500	6.08333	16.00000	3.32500	43,815.14410	0.00000	
4	0.00000	90.16667	6.02500	16.00000	3.61667	43,815.18576	0.00000	
...	...	...	...	...	...	...	...	
3979	0.00000	40.16667	20.14167	2.00000	4.20000	43,980.81076	0.00000	
3980	0.00000	44.49167	19.29167	2.00000	3.99167	43,980.85243	0.00000	
3981	0.00000	62.33333	17.42500	2.00000	4.07500	43,980.89410	0.00000	
3982	0.00000	68.09167	16.07500	2.00000	4.30000	43,980.93576	0.00000	
3983	0.00000	71.32500	15.29167	2.00000	4.21667	43,980.97743	0.00000	

	BlackScr	CO2air	Cum_irr	EC_drain_PC	EnScr	HumDef	PipeGrow	\
0	95.00000	493.36364	0.53220	3.98909	100.00000	6.89000	0.00000	
1	96.00000	498.66667	0.53220	3.97583	100.00000	6.95917	0.00000	
2	96.00000	494.91667	0.04435	3.97417	100.00000	7.04167	0.00000	
3	96.00000	494.00000	0.00000	3.96750	100.00000	7.03167	0.00000	
4	96.00000	500.00000	0.00000	3.96917	100.00000	7.02000	0.00000	
...	...	...	...	...	...	...	...	
3979	0.00000	424.50000	1.59660	10.17083	0.00000	12.19417	0.00000	
3980	0.00000	431.33333	1.59660	10.19667	0.00000	10.38667	0.00000	
3981	0.00000	435.66667	1.59660	10.26417	0.00000	6.93167	0.00000	
3982	0.00000	449.54167	1.59660	10.29583	0.00000	4.89917	0.00000	
3983	0.00000	464.00000	1.59660	10.32750	0.83333	3.05000	0.00000	

	PipeLow	Rhair	Tair	Tot_PAR	Tot_PAR_Lamps	VentLee	Ventwind	\
0	43.03636	61.82727	20.76364	0.00000	0.00000	0.80000	0.00000	

1	46.88333	61.39167	20.74167	0.00000	0.00000	0.80000	0.00000
2	46.84167	61.33333	20.92500	0.00000	0.00000	0.80000	0.00000
3	47.13333	61.57500	21.00833	0.00000	0.00000	0.80000	0.00000
4	47.37500	61.88333	21.11667	0.00000	0.00000	0.80000	0.00000
...	...	...	...	...	...	...	...
3979	0.00000	40.37500	22.91667	263.41667	0.00000	100.00000	78.05000
3980	0.00000	44.72500	21.44167	117.66667	0.00000	100.00000	57.86667
3981	0.00000	59.48333	19.80000	22.45833	0.00000	87.56667	8.66667
3982	0.00000	69.52500	18.77917	0.50000	0.00000	39.54167	0.00000
3983	0.00000	80.54167	18.35833	0.16667	0.00000	0.00000	0.00000

	assim_vip	co2_dos	co2_vip	dx_vip	int_farred_sp	int_white_sp	\
0	0.00000	0.00078	400.00000	2.20000	0.00000	0.00000	
1	0.00000	0.00034	400.00000	2.20000	0.00000	0.00000	
2	0.00000	0.00006	400.00000	2.20000	0.00000	0.00000	
3	0.00000	0.00012	400.00000	2.20000	0.00000	0.00000	
4	0.00000	0.00018	400.00000	2.20000	0.00000	0.00000	
...	...	...	...	...	...	...	
3979	0.00000	0.00125	400.00000	2.20000	0.00000	0.00000	
3980	0.00000	0.00113	400.00000	2.20000	0.00000	0.00000	
3981	0.00000	0.00088	400.00000	2.20000	0.00000	0.00000	
3982	0.00000	0.00056	400.00000	2.20000	0.00000	0.00000	
3983	0.00000	0.00023	400.00000	2.20000	0.00000	0.00000	

	pH_drain_PC	scr_blk_vip	scr_enrg_vip	t_grow_min_vip	t_heat_vip	\
0	6.27909	96.00000	100.00000	10.00000	21.00000	
1	6.27917	96.00000	100.00000	10.00000	21.00000	
2	6.27833	96.00000	100.00000	10.00000	21.00000	
3	6.27667	96.00000	100.00000	10.00000	21.00000	
4	6.28000	96.00000	100.00000	10.00000	21.00000	
...	...	...	...	...	...	
3979	4.42833	96.00000	100.00000	0.00000	10.00000	
3980	4.44000	96.00000	85.00000	0.00000	10.00000	
3981	4.45250	96.00000	80.00000	0.00000	10.00000	
3982	4.46833	96.00000	98.33333	0.00000	10.00000	
3983	4.47583	96.00000	100.00000	0.00000	10.00000	

	t_rail_min_vip	t_ventlee_vip	t_ventwind_vip	water_sup	\
0	0.00000	25.00000	26.00000	6.00000	
1	0.00000	25.00000	26.00000	6.00000	
2	0.00000	25.00000	26.00000	0.50000	
3	0.00000	25.00000	26.00000	0.00000	
4	0.00000	25.00000	26.00000	0.00000	
...	...	...	...	...	
3979	0.00000	18.00000	20.00000	18.00000	
3980	0.00000	18.00000	20.00000	18.00000	
3981	0.00000	18.00000	20.00000	18.00000	

3982	0.00000	18.00000	20.00000	18.00000
3983	0.00000	18.00000	20.00000	18.00000

	water_sup_intervals_vip_min	window_pos_lee_vip
0	1,000.00000	1.20000
1	1,000.00000	1.20000
2	1,000.00000	1.20000
3	1,000.00000	1.20000
4	1,000.00000	1.20000
...	...	...
3979	30.00000	0.00000
3980	30.00000	0.00000
3981	30.00000	0.00000
3982	30.00000	0.00000
3983	30.00000	0.00000

[3983 rows x 44 columns]

```
[87]: start_date = '2013-01-01'
end_date = '2023-11-27'
step = pd.DateOffset(days=1)

df_combined_hr_index = pd.date_range(start=start_date, end=end_date, freq=step)
df_combined_hr.index = df_combined_hr_index
```

```
[ ]:
```

```
[88]: from sklearn.feature_selection import SelectKBest, f_classif

target_variable = df_combined_hr['Tair']

# Selecting the features using SelectKBest and f_classif
selector = SelectKBest(f_classif, k=23)
X_new = selector.fit_transform(df_combined_hr, target_variable)
names = df_combined_hr.columns.values[selector.get_support()]
scores = selector.scores_[selector.get_support()]
names_scores = list(zip(names, scores))
ns_df = pd.DataFrame(data=names_scores, columns=['Feat_names', 'F_Scores'])
ns_df_sorted = ns_df.sort_values(['F_Scores', 'Feat_names'], ascending = [
    ↪ [False, True])
print(ns_df_sorted)
```

	Feat_names	F_Scores
9	Tair	inf
15	dx_vip	inf
18	t_heat_vip	15.86973
19	t_ventlee_vip	11.24211

20	t_ventwind_vip	9.18111
10	Tot_PAR	4.62515
17	scr_enrg_vip	3.63696
12	Ventwind	3.45646
4	BlackScr	2.99139
5	Cum_irr	2.79312
21	water_sup	2.79312
1	PARout	2.68716
0	Iglob	2.65395
6	HumDef	2.50714
13	assim_vip	2.13078
22	water_sup_intervals_vip_min	2.12508
3	AssimLight	2.12007
11	VentLee	1.99755
7	PipeGrow	1.94458
16	int_farred_sp	1.86257
8	Rhair	1.83319
2	RadSum	1.75383
14	co2_vip	1.74156

```
[89]: print ("Most influential variables: dx_vip: Humidity deficit VIP (Vapor_
↳Pressure Deficit ), \
t_ventlee_vip: Ventilation temperature VIP (leeward vents, temperature of the_
↳incoming air) \
t_ventwind_vip: Ventilation temperature VIP (windward side), t_heat_vip:_
↳Heating temperature VIP")
```

Most influential variables: dx\_vip: Humidity deficit VIP (Vapor Pressure Deficit ), t\_ventlee\_vip: Ventilation temperature VIP (leeward vents, temperature of the incoming air) t\_ventwind\_vip: Ventilation temperature VIP (windward side), t\_heat\_vip: Heating temperature VIP

```
[90]: print ("The dx_vip is directly related to the temperature, and actually is the_
↳temperature \
that influences the dx_vip. So I won't use this variable for the predition of_
↳Tair.")
```

The dx\_vip is directly related to the temperature, and actually is the temperature that influences the dx\_vip. So I won't use this variable for the predition of Tair.

```
[91]: combined_train = df_combined_hr.iloc[-500:-24]

# Create the test set
combined_test = df_combined_hr.iloc[-24:]
```

```
[92]: from statsmodels.tsa.statespace.sarimax import SARIMAX
endog = combined_train['Tair']

# Define exogenous variables
exog_train = combined_train[['t_ventlee_vip', 't_ventwind_vip', 't_heat_vip']]

# Set the start index for prediction
start_date = combined_train.index[-1] + timedelta(days=1) # Assuming your data
↳ has hourly frequency

# Set the end index for prediction
end_date = start_date + timedelta(days=24)

# Specify the SARIMAX model with exogenous variables
sarimax_model = SARIMAX(endog, exog=exog_train, order=(1, 0, 1),
↳ seasonal_order=(2, 1, 0, 24))

# Fit the model
results = sarimax_model.fit()

# Obtain summary statistics
summary = results.summary()


# Create exogenous variables for prediction
exog_pred = exog_train[-25:] # Get the last 25 rows of exogenous variables

# Make predictions with confidence intervals
predictions = results.get_prediction(start=start_date, end=end_date,
↳ exog=exog_pred, dynamic=False)
predicted_values = predictions.predicted_mean
confidence_intervals = predictions.conf_int()
lower_percentile = np.percentile(confidence_intervals, 10, axis=1)
upper_percentile = np.percentile(confidence_intervals, 90, axis=1)

# Plotting the predictions with confidence intervals
plt.figure(figsize=(15, 7))
plt.plot(combined_train['Tair'].tail(50), label='Actual')
plt.plot(predicted_values, color='red', label='Predicted')
plt.plot(combined_test['Tair'], color='black', label='Actual')
# plt.fill_between(confidence_intervals.index, confidence_intervals.iloc[:, 0],
↳ confidence_intervals.iloc[:, 1], color='gray', alpha=0.3)
plt.fill_between(confidence_intervals.index, lower_percentile,
↳ upper_percentile, color='k', alpha=.15, label='90% Confidence Interval')

plt.title("SARIMAX - Forecast with Confidence Intervals")
plt.legend()
```

```
plt.show()
```



```
[93]: print ("summary statistics")
      print (summary)
```

summary statistics

#### SARIMAX Results

```
=====
=====
Dep. Variable:          Tair    No. Observations:
476
Model:                SARIMAX(1, 0, 1)x(2, 1, [], 24)    Log Likelihood
-362.862
Date:                  Sun, 09 Jul 2023    AIC
741.723
Time:                  14:51:03    BIC
774.633
Sample:                07-16-2022    HQIC
754.692
- 11-03-2023
```

Covariance Type:

opg

```
=====
==
coef    std err        z    P>|z|    [0.025
0.975]
-----
--
t_ventlee_vip    -0.2946    0.423    -0.697    0.486    -1.123
0.533
```

t_ventwind_vip	0.4813	0.069	6.963	0.000	0.346
0.617					
t_heat_vip	0.6565	0.452	1.451	0.147	-0.230
1.543					
ar.L1	0.8485	0.020	41.974	0.000	0.809
0.888					
ma.L1	0.2613	0.047	5.571	0.000	0.169
0.353					
ar.S.L24	-0.4731	0.037	-12.897	0.000	-0.545
-0.401					
ar.S.L48	-0.3199	0.045	-7.183	0.000	-0.407
-0.233					
sigma2	0.2851	0.014	20.639	0.000	0.258
0.312					

=====

===

Ljung-Box (L1) (Q):	0.08	Jarque-Bera (JB):
99.87		

Prob(Q):	0.78	Prob(JB):
0.00		

Heteroskedasticity (H):	1.33	Skew:
-0.14		

Prob(H) (two-sided):	0.08	Kurtosis:
5.29		

=====

===

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
[94]: from sklearn.metrics import mean_absolute_error, mean_squared_error

# Define the actual values
actual_values = combined_test['Tair']

# Calculate the MAE
maeSX1 = mean_absolute_error(actual_values, predicted_values[-24:])

# Calculate the MSE
mseSX1 = mean_squared_error(actual_values, predicted_values[-24:])

# Calculate the RMSE (using mseSX1 instead of mse)
rmseSX1 = np.sqrt(mseSX1)

# Calculate the MAPE
```



```
mapeSX1 = np.mean(np.abs((actual_values - predicted_values[-24:]) /
↳actual_values)) * 100

# Print the metrics
print("MAE:", maeSX1)
print("MSE:", mseSX1)
print("RMSE:", rmseSX1)
print("MAPE:", mapeSX1)
```

```
MAE: 0.8501748698533372
MSE: 1.055535696933144
RMSE: 1.0273926693008588
MAPE: 4.356218604685104
```

```
[95]: comment = "**7. Final Comparison of Models Using MSE**"
font_size = "20px"

Markdown(f"<span style='font-weight:bold; font-size:{font_size}'>{comment}</
↳span>")
```

[95]: 7. Final Comparison of Models Using MSE

```
[96]: print("FINAL COMPARISON")
data_dict = {'MSE':[np.mean(simple_mse), np.mean(triple_mse), maeSA, maeSX1]}
df = pd.DataFrame(data_dict, index=['simple averaged smoothing','triple
↳exponential smoothing'\
                                     , 'SARIMA', 'SARIMAX'])
print(df)
```

```
FINAL COMPARISON
                                     MSE
simple averaged smoothing    470.80015
triple exponential smoothing 682.13334
SARIMA                      0.59406
SARIMAX                      0.85017
```

[ ]: